

Doc2Vec, UMAP and HDBSCAN

May 21, 2020

Uses:

<http://linanqiu.github.io/2015/10/07/word2vec-sentiment/>

<https://umap-learn.readthedocs.io/en/latest/clustering.html>

0.0.1 Train Doc2Vec

```
In [1]: # gensim modules
        from gensim import utils
        from gensim.models.doc2vec import TaggedDocument
        from gensim.models import Doc2Vec

        # numpy
        import numpy

        # random
        from random import shuffle

        # classifier
        from sklearn.linear_model import LogisticRegression

In [ ]: # got data by cloning: https://github.com/linanqiu/word2vec-sentiments
        #gensim Doc2Vec only takes in Labeled Line Sentence classes

In [2]: #class to take a dictionary of files, read, format, label
        #classes are objects of functions that share variables

        class LabeledLineSentence(object):
            #initialize class. self refers to the class
            def __init__(self, sources):
                #sources will be a shared parameter for class
                self.sources = sources

            flipped = {}

            # make sure that keys are unique (sources is a dict)
            #flip key:value to value:key to ensure every key has unique value
            #raise exception if have multiple keys for 1 value
            for key, value in sources.items():
```

```

        if value not in flipped:
            flipped[value] = [key]
        else:
            raise Exception('Non-unique prefix encountered')

#if want data as iterator
#iterate through sources
#for each source, open file
#return unicode list of words in line split by space and prefix+item# as TaggedDoc
def __iter__(self):
    for source, prefix in self.sources.items():
        with utils.smart_open(source) as fin:
            for item_no, line in enumerate(fin):
                yield TaggedDocument(utils.to_unicode(line).split(), [prefix + '_%s' % item_no])

#if want data as an array
def to_array(self):
    self.sentences = []
    for source, prefix in self.sources.items():
        with utils.smart_open(source) as fin:
            for item_no, line in enumerate(fin):
                self.sentences.append(TaggedDocument(utils.to_unicode(line).split(), [prefix + '_%s' % item_no]))
    return self.sentences

#shuffle document order
def sentences_perm(self):
    shuffled = list(self.sentences)
    shuffle(shuffled)
    return shuffled

In [3]: #create dict of data sources
sources = {'w2vsent (git)/word2vec-sentiments/test-neg.txt': 'TEST_NEG', 'w2vsent (git)/word2vec-sentiments/test-pos.txt': 'TEST_POS'}

sentences = LabeledLineSentence(sources)

In [18]: #define doc2vec model
model = Doc2Vec(min_count=1, window=10, vector_size=100, sample=1e-4, negative=5, workers=-1)

#use model to build vocab
model.build_vocab(sentences.to_array())

In [21]: #multiple epochs, but shuffle the order we feed in after each epoch
for epoch in range(10):
    print('Epoch %d' % epoch)
    model.train(sentences.sentences_perm(), total_examples=model.corpus_count, epochs=1)

Epoch 0
Epoch 1

```

Epoch 2
Epoch 3
Epoch 4
Epoch 5
Epoch 6
Epoch 7
Epoch 8
Epoch 9

```
In [23]: model.wv.most_similar('good')
```

```
Out[23]: [('nice', 0.6974087953567505),  
          ('great', 0.6966586709022522),  
          ('bad', 0.6601657271385193),  
          ('decent', 0.6535385847091675),  
          ('terrific', 0.612062931060791),  
          ('solid', 0.6077233552932739),  
          ('fine', 0.5853490829467773),  
          ('excellent', 0.5782768726348877),  
          ('alright', 0.559374213218689),  
          ('fantastic', 0.5098010301589966)]
```

```
In [24]: model.wv.most_similar('spielberg')
```

```
Out[24]: [('steven', 0.5286897420883179),  
          ('serpico', 0.5100363492965698),  
          ('latt', 0.5079025030136108),  
          ('bashers', 0.5069671273231506),  
          ('kubrick', 0.5027161836624146),  
          ('soderberg', 0.4896114766597748),  
          ('craven', 0.4726443886756897),  
          ('ruegger', 0.4686766266822815),  
          ('gulagher', 0.4619401693344116),  
          ('coppola', 0.460299015045166)]
```

```
In [28]: model['TRAIN_POS_0']
```

```
Out[28]: array([ 0.17210831,  0.53068894,  0.08875216, -0.22204731,  0.08920082,  
                -0.41164    , -0.10716269,  0.01384323, -0.44027948, -0.04258307,  
                -0.3661883 ,  0.3593273 , -0.39790803,  0.29294032, -0.5823881 ,  
                0.4351405 , -0.33027872,  0.05932414,  0.50758123, -0.08620853,  
                0.03721116, -0.0885262 ,  0.47685373,  0.0458203 , -0.1292882 ,  
                0.32321823,  0.1800003 , -0.3754428 , -0.10352938,  0.02824782,  
                0.50372    ,  0.13023265, -0.11537365,  0.10333041, -0.5935597 ,  
                0.10701659,  0.42890313, -0.4348067 , -0.05726193, -0.42349628,  
                -0.14245322,  0.15413284,  1.047601  ,  0.13994423,  0.8603275 ,  
                -0.30993137,  0.11837404, -0.04155462,  0.18748598, -0.6198913 ,  
                -0.1445792 ,  0.00320736, -0.37068844,  0.32822382, -0.17440598,
```

```

-0.24452095, 0.19037859, -0.43866974, -0.8702648 , -0.7223541 ,
0.21947137, 0.42983302, -0.05602715, 0.17835127, 0.06716342,
-0.31129426, -0.20862402, -0.5284197 , -0.36720476, -0.31597954,
-0.11324366, -0.53116 , -0.03544942, 0.31802335, 0.52840805,
0.39634323, -0.5538435 , 0.17646116, 0.20288238, -0.31766 ,
0.15241203, 0.1531507 , -0.556646 , 0.3208777 , 0.01118903,
-0.0276349 , -0.6513774 , 0.03072487, -0.28013617, 0.15971902,
-0.13761188, 0.03399257, 0.14500739, -0.28899762, 0.23199414,
0.33987027, 0.47027183, 0.34716403, 0.16056226, -0.13079283],
dtype=float32)

```

```
In [87]: model.docvecs.most_similar('TRAIN_POS_0')
```

```

Out[87]: [('TRAIN_POS_1934', 0.6886619329452515),
('TRAIN_POS_4722', 0.6827501058578491),
('TEST_POS_6954', 0.6789910793304443),
('TRAIN_POS_8849', 0.6780802011489868),
('TRAIN_NEG_3214', 0.6722345352172852),
('TRAIN_NEG_4430', 0.6705719232559204),
('TRAIN_POS_1496', 0.668876051902771),
('TEST_POS_7943', 0.6654232740402222),
('TRAIN_NEG_7613', 0.6619976162910461),
('TRAIN_POS_9979', 0.6605291962623596)]

```

0.0.2 UMAP

UMAP as preprocessing useful for density based clustering

```

In [95]: from sklearn.decomposition import PCA
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# Dimension reduction and clustering libraries
import umap
import hdbscan
from sklearn.metrics import adjusted_rand_score, adjusted_mutual_info_score

```

```
In [142]: sns.set(style='white', rc={'figure.figsize':(10,8)})
```

```
In [90]: #extract document vectors and labels
```

```

train_arrays = numpy.zeros((25000, 100))
train_labels = numpy.zeros(25000)

for i in range(12500):
    prefix_train_pos = 'TRAIN_POS_' + str(i)

```

```

prefix_train_neg = 'TRAIN_NEG_' + str(i)
train_arrays[i] = model[prefix_train_pos]
train_arrays[12500 + i] = model[prefix_train_neg]
train_labels[i] = 1
train_labels[12500 + i] = 0


```

In [93]: *#umap and viz colored by positive/negative "labels"*

```

standard_embedding = umap.UMAP(random_state=42).fit_transform(train_arrays)
plt.scatter(standard_embedding[:, 0], standard_embedding[:, 1], c=train_labels, s=0.1

```



0.0.3 HDBSCAN

extends DBSCAN by converting into hierarchical clustering algorithm and then extracting a flat clustering based on the stability of clusters. - Transforms space according to density/sparsity - Build minimum spanning tree of the distance weighted graph - Constructs a cluster hierarchy based on min cluster size - Extracts the stable clusters from the condensed tree

In [96]: *# reduce dimensions to 50 with PCA*

```

lowd_imdb = PCA(n_components=50).fit_transform(train_arrays)

```

#run HDBSCAN

```


hdbscan_labels = hdbscan.HDBSCAN(min_samples=10, min_cluster_size=500).fit_predict(lowd_imdb)

```

```

In [97]: #hdbscan defers clustering points it isn't confident in,
#classifying them as noise
#check this by coloring points not clasified gray
clustered = (hdbscan_labels >= 0)
plt.scatter(standard_embedding[~clustered, 0],
            standard_embedding[~clustered, 1],
            c=(0.5, 0.5, 0.5),
            s=0.1,
            alpha=0.5)
plt.scatter(standard_embedding[clustered, 0],
            standard_embedding[clustered, 1],
            c=hdbscan_labels[clustered],
            s=0.1,
            cmap='Spectral');

```



```

In [101]: #it actually didn't cluster anything
np.sum(clustered)

#problem with high dimensions in density clustering algoritms.
#high dimensions requires more observed samples to produce density
#However reducing any more with PCA will lose a lot of variance in the data

```

Out[101]: 0

```

In [137]: # when using UMAP for dimension reduction instead of visualization,
# want larger n_neighbors value to maintain more global structure

```

```

# want low min_dist because want to pack points densely

clusterable_embedding = umap.UMAP(
    n_neighbors=30,
    min_dist=0.0,
    n_components=3,
    random_state=42,
).fit_transform(train_arrays)


In [157]: hdb_clusterer=hdbscan.HDBSCAN(
    min_samples=5,
    min_cluster_size=250)

labels = hdb_clusterer.fit_predict(clusterable_embedding)

In [158]: #check how many points clustered

clustered = (labels >= 0)
plt.scatter(standard_embedding[~clustered, 0],
            standard_embedding[~clustered, 1],
            c=(0.5, 0.5, 0.5),
            s=0.1,
            alpha=0.5)
plt.scatter(standard_embedding[clustered, 0],
            standard_embedding[clustered, 1],
            c=labels[clustered],
            s=0.1,
            cmap='Spectral');

```



```
In [159]: np.sum(clustered)/train_arrays.shape[0]
```

```
Out[159]: 0.54548
```


```
In [160]: #how many clusters?  
          set(labels)
```

```
Out[160]: {-1, 0, 1, 2, 3, 4}
```

```
In [161]: #dendrogram
```

```
          hdb_clusterer.condensed_tree_.plot(select_clusters=True, selection_palette=sns.color
```

```
Out[161]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1ce02c88>
```

In []: