

UMAP Presentation Example

May 21, 2020

1 UMAP

- Iris vs TSNE
- MNIST vs TSNE
- Parameter selection on generated data
- Comparison to other dimensionality reduction algorithms
- Word Vectors from NIPS conference paper abstracts

Installation: <https://umap-learn.readthedocs.io/en/latest/index.html>

IRIS and MNIST sections follow: https://umap-learn.readthedocs.io/en/latest/basic_usage.html

```
In [69]: import umap
         from sklearn.manifold import TSNE
         import numpy as np
         import matplotlib.pyplot as plt
         from mpl_toolkits.mplot3d import Axes3D
         import seaborn as sns
         import pandas as pd
         import time
         from sklearn import datasets, decomposition, manifold, preprocessing
         from colorsys import hsv_to_rgb
         import re
         from nltk.corpus import stopwords
         from nltk.tokenize import RegexpTokenizer
         from gensim.models import word2vec
         %matplotlib inline

         #Load iris and digits datasets
         from sklearn.datasets import load_iris, load_digits

         import warnings
         warnings.filterwarnings('ignore') # action='once'

In [70]: sns.set(style='white', context='notebook', rc={'figure.figsize':(12,8)})
```

1.0.1 Iris dataset

```
In [5]: iris = load_iris()
```

```
#4 dimensions - sepal length, sepal width, petal length, petal width  
#see it as 2d scatterplots
```

```
iris_df = pd.DataFrame(iris.data, columns=iris.feature_names)  
iris_df['species'] = pd.Series(iris.target).map(dict(zip(range(3), iris.target_names)))
```

```
In [24]: #UMAP follows sklearn API
```

```
reducer = umap.UMAP()
```

```
%time umap_iris = reducer.fit_transform(iris.data)  
print(umap_iris.shape)  
umap_iris[:5]
```


```
CPU times: user 335 ms, sys: 35.8 ms, total: 370 ms
```

```
Wall time: 345 ms
```

```
(150, 2)
```

```
Out[24]: array([[ 8.878182 , -3.778602 ],  
                [10.777144 , -4.5306053],  
                [10.74014 , -3.9289017],  
                [10.774868 , -3.9556265],  
                [ 9.01998 , -3.8093293]], dtype=float32)
```

```
In [8]: plt.scatter(umap_iris[:, 0], umap_iris[:, 1], c=[sns.color_palette()[x] for x in iris.target_names])  
plt.gca().set_aspect('equal', 'datalim')  
plt.title('UMAP projection of the Iris dataset', fontsize=24);
```



```
In [23]: reducer = TSNE()
```

```
tsne_iris = reducer.fit_transform(iris.data)
print(tsne_iris.shape)
tsne_iris[:5]
```

```
%time umap_iris = reducer.fit_transform(iris.data)
print(umap_iris.shape)
print(umap_iris[:5])
```

```
(150, 2)
```

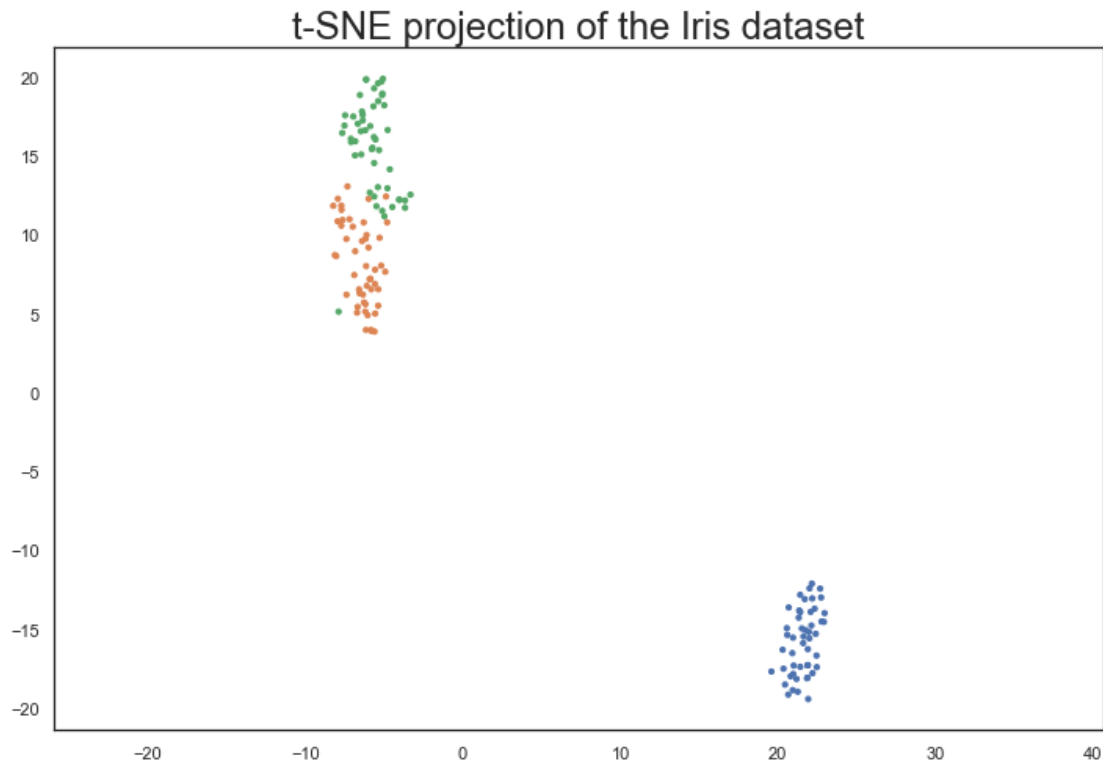
```
CPU times: user 2.97 s, sys: 336 ms, total: 3.3 s
```

```
Wall time: 3.31 s
```

```
(150, 2)
```

```
[[-19.42122    9.890934]
 [-16.871984  11.127105]
 [-16.98465    9.918852]
 [-16.546099  10.34091 ]
 [-19.469473   9.630399]]
```

```
In [13]: plt.scatter(tsne_iris[:, 0], tsne_iris[:, 1], c=[sns.color_palette()[x] for x in iris
plt.gca().set_aspect('equal', 'datalim')
plt.title('t-SNE projection of the Iris dataset', fontsize=24);
```



1.0.2 MNIST dataset

```
In [19]: digits = load_digits()
```

```
#8x8 images so 64 attributes
```

```
#examples
```

```
fig, ax_array = plt.subplots(15, 15)
```

```
axes = ax_array.flatten()
```

```
for i, ax in enumerate(axes):
```

```
    ax.imshow(digits.images[i], cmap='gray_r')
```

```
plt.setp(axes, xticks=[], yticks=[], frame_on=False)
```

```
plt.tight_layout(h_pad=0.5, w_pad=0.01)
```

0	1	2	3	4	5	6	7	8	9	0	1	2	3	4
5	6	7	8	9	0	1	2	3	4	5	6	7	8	9
0	3	5	5	6	5	0	3	8	9	8	4	1	7	7
3	5	1	0	0	2	2	7	8	2	0	1	2	6	3
3	7	3	3	4	6	6	6	4	7	1	5	0	3	5
2	8	2	0	0	1	7	6	3	2	1	7	4	6	3
1	3	3	1	7	6	8	4	3	1	4	0	5	3	6
3	6	1	7	5	4	4	7	2	8	2	2	5	7	9
5	4	8	8	4	9	0	1	9	1	0	1	2	3	4
5	6	7	8	9	0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9	0	3	5	5	6
5	0	9	8	9	8	4	1	7	7	3	5	1	0	0
2	2	7	8	1	0	1	2	6	3	3	7	3	3	4
6	6	6	4	9	1	5	0	9	5	2	8	1	0	0
1	7	6	3	2	1	7	3	1	3	9	1	7	6	8

```
In [38]: reducer = umap.UMAP(random_state=42)
         %time reducer.fit(digits.data)
```

```
CPU times: user 4.27 s, sys: 348 ms, total: 4.62 s
Wall time: 4.45 s
```

```
Out[38]: UMAP(a=None, angular_rp_forest=False, b=None, init='spectral',
             learning_rate=1.0, local_connectivity=1.0, metric='euclidean',
             metric_kws=None, min_dist=0.1, n_components=2, n_epochs=None,
             n_neighbors=15, negative_sample_rate=5, random_state=42,
             repulsion_strength=1.0, set_op_mix_ratio=1.0, spread=1.0,
             target_metric='categorical', target_metric_kws=None,
             target_n_neighbors=-1, target_weight=0.5, transform_queue_size=4.0,
             transform_seed=42, verbose=False)
```


```
In [17]: umap_digit = reducer.transform(digits.data)

         print(umap_digit.shape)
         umap_digit[:5]
```

```
(1797, 2)
```

```
Out[17]: array([[15.318228 ,  1.892302 ],
                [-7.0599804,  2.852144 ],
                [-5.364305 ,  1.1695055 ],
                [-0.96996576,  0.36345217],
                [ 6.3392377 , 10.887402 ]], dtype=float32)
```

```
In [18]: plt.scatter(umap_digit[:, 0], umap_digit[:, 1], c=digits.target, cmap='Spectral', s=50,
                    plt.gca().set_aspect('equal', 'datalim')
                    plt.colorbar(boundaries=np.arange(11)-0.5).set_ticks(np.arange(10))
                    plt.title('UMAP projection of the Digits dataset', fontsize=24);
```




```
In [27]: reducer = TSNE(random_state=42)
         %time tsne_digit = reducer.fit_transform(digits.data)
         print(tsne_digit.shape)
         tsne_digit[:5]
```

```
CPU times: user 46.5 s, sys: 5.53 s, total: 52 s
Wall time: 53.1 s
(1797, 2)
```

```
Out[27]: array([[ 64.89154   , -0.91041076],
                [-23.001663 ,  0.37934512],
                [ -7.88827   , 27.260277  ],
                [ 26.571157  , 17.319496  ],
                [-43.865337 , -27.682026  ]], dtype=float32)
```

```
In [28]: plt.scatter(tsne_digit[:, 0], tsne_digit[:, 1], c=digits.target, cmap='Spectral', s=50,
                    plt.gca().set_aspect('equal', 'datalim')
                    plt.colorbar(boundaries=np.arange(11)-0.5).set_ticks(np.arange(10))
                    plt.title('t-SNE projection of the Digits dataset', fontsize=24);
```



1.0.3 Random 4D cube

The major UMAP parameters are:

- `n_neighbors` - determines balance of local vs global structure. It constrains the size of the local neighborhood UMAP looks at when attempting to learn manifold structure. Low values brings out local structure, large bring out larger neighborhoods while losing the fine detail.
- `min_dist` - determines the minimum distance apart points are allowed to be in the embedding. Low values will result in clumpier embeddings which is good for finer topological structure. Larger values is useful to visualize broad topological structure.

- `n_components` - how many components to return. t-SNE isn't great in more than 2 dimensions but UMAP is.
- `metric` - controls how the distance is computed in input data. Many metrics are built in but they can also be defined if compiled with numba.

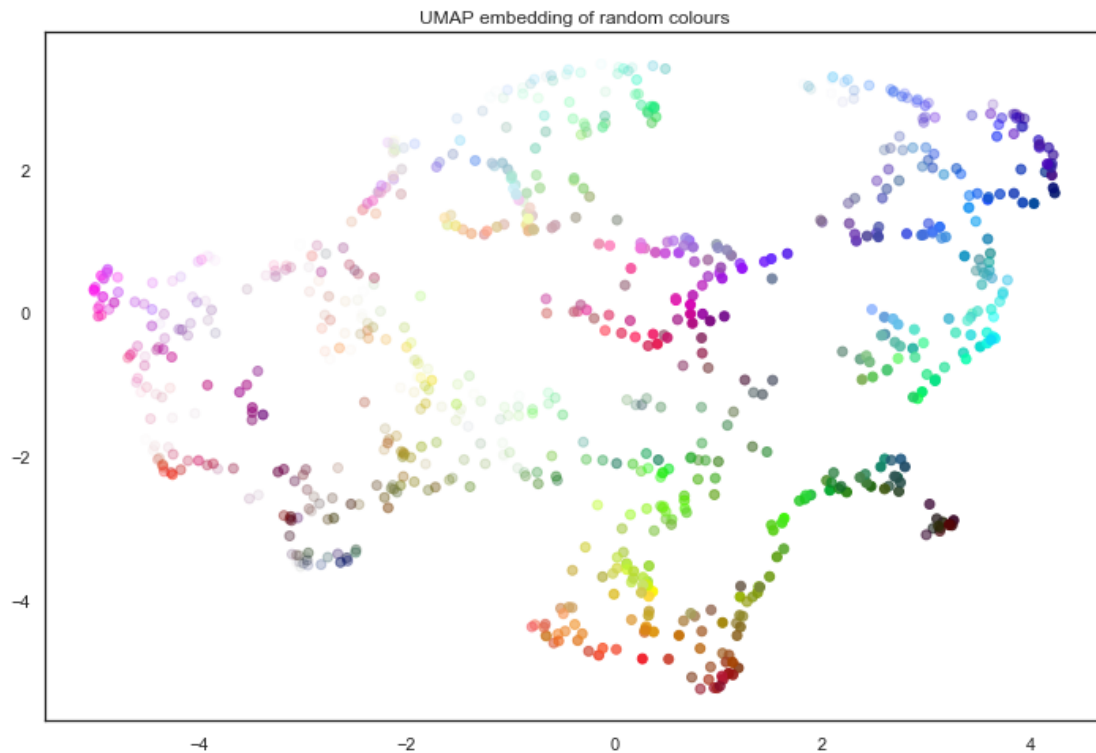
```
In [66]: #generate 4d cube
np.random.seed(42)
data = np.random.rand(800, 4)
```

```
In [67]: fit = umap.UMAP()
%time u = fit.fit_transform(data)
```

CPU times: user 1.67 s, sys: 202 ms, total: 1.88 s

Wall time: 1.88 s

```
In [71]: #colors represent the data. So data close to one another in 4d space are
#similar colors
plt.scatter(u[:,0], u[:,1], c=data)
plt.title('UMAP embedding of random colors');
```



```
In [72]: #function which fits, transforms and plots
```

```
def draw_umap(n_neighbors=15, min_dist=0.1, n_components=2, metric='euclidean', title=
```



```

fit = umap.UMAP(
    n_neighbors=n_neighbors,
    min_dist=min_dist,
    n_components=n_components,
    metric=metric
)
u = fit.fit_transform(data);
fig = plt.figure()
if n_components == 1:
    ax = fig.add_subplot(111)
    ax.scatter(u[:,0], range(len(u)), c=data)
if n_components == 2:
    ax = fig.add_subplot(111)
    ax.scatter(u[:,0], u[:,1], c=data)
if n_components == 3:
    ax = fig.add_subplot(111, projection='3d')
    ax.scatter(u[:,0], u[:,1], u[:,2], c=data, s=100)
plt.title(title, fontsize=18)

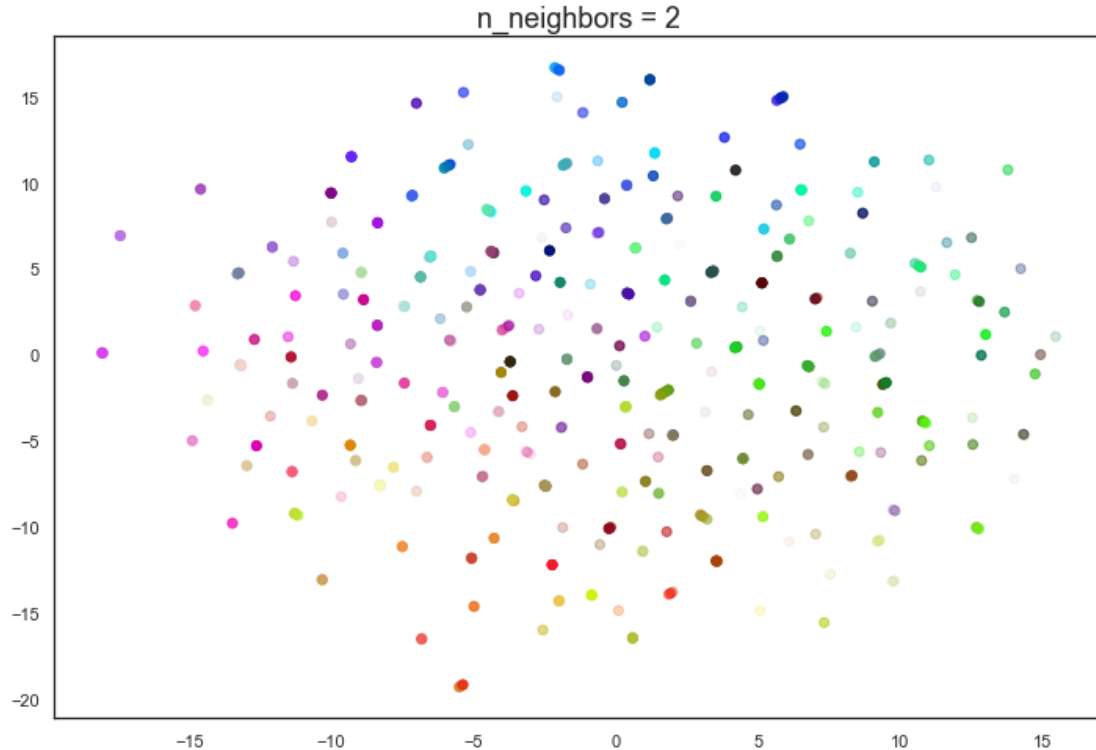
```



```



In [73]: for n in (2, 5, 10, 20, 50, 100, 200):
        draw_umap(n_neighbors=n, title='n_neighbors = {}'.format(n))


```


#can see that higher n_neighbors glues more points together

















```
In [74]: for d in (0.0, 0.1, 0.25, 0.5, 0.8, 0.99):  
         draw_umap(min_dist=d, title='min_dist = {}'.format(d))  
  
         # clumps and strings pushed apart as min_dist is greater
```

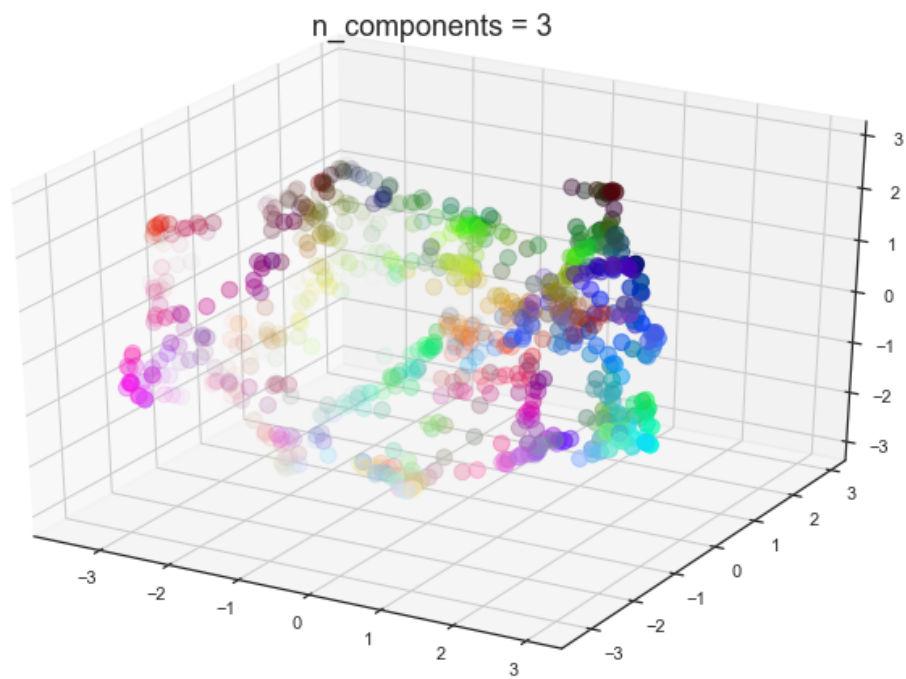








```
In [75]: draw_umap(n_components=3, title='n_components = 3')
```



1.0.4 Comparison of Dimensionality Techniques

Taken directly from: https://umap-learn.readthedocs.io/en/latest/auto_examples/plot_algorithm_comparison_glr-auto-examples-plot-algorithm-comparison-py

```
In [30]: sns.set(context="paper", style="white")

blobs, blob_labels = datasets.make_blobs(
    n_samples=500, n_features=10, centers=5, random_state=42
)
iris = datasets.load_iris()
digits = datasets.load_digits(n_class=10)
wine = datasets.load_wine()
swissroll, swissroll_labels = datasets.make_swiss_roll(
    n_samples=1000, noise=0.1, random_state=42
)
sphere = np.random.normal(size=(600, 3))
sphere = preprocessing.normalize(sphere)
sphere_hsv = np.array(
    [
        (
            (np.arctan2(c[1], c[0]) + np.pi) / (2 * np.pi),
            np.abs(c[2]),
            min((c[2] + 1.1), 1.0),
        )
        for c in sphere
    ]
)
sphere_colors = np.array([hsv_to_rgb(*c) for c in sphere_hsv])

reducers = [
    (manifold.TSNE, {"perplexity": 50}),
    # (manifold.LocallyLinearEmbedding, {'n_neighbors':10, 'method':'hessian'}),
    (manifold.Isomap, {"n_neighbors": 30}),
    (manifold.MDS, {}),
    (decomposition.PCA, {}),
    (umap.UMAP, {"n_neighbors": 30, "min_dist": 0.3}),
]

test_data = [
    (blobs, blob_labels),
    (iris.data, iris.target),
    (digits.data, digits.target),
    (wine.data, wine.target),
    (swissroll, swissroll_labels),
```

```

        (sphere, sphere_colors),
    ]
    dataset_names = ["Blobs", "Iris", "Digits", "Wine", "Swiss Roll", "Sphere"]

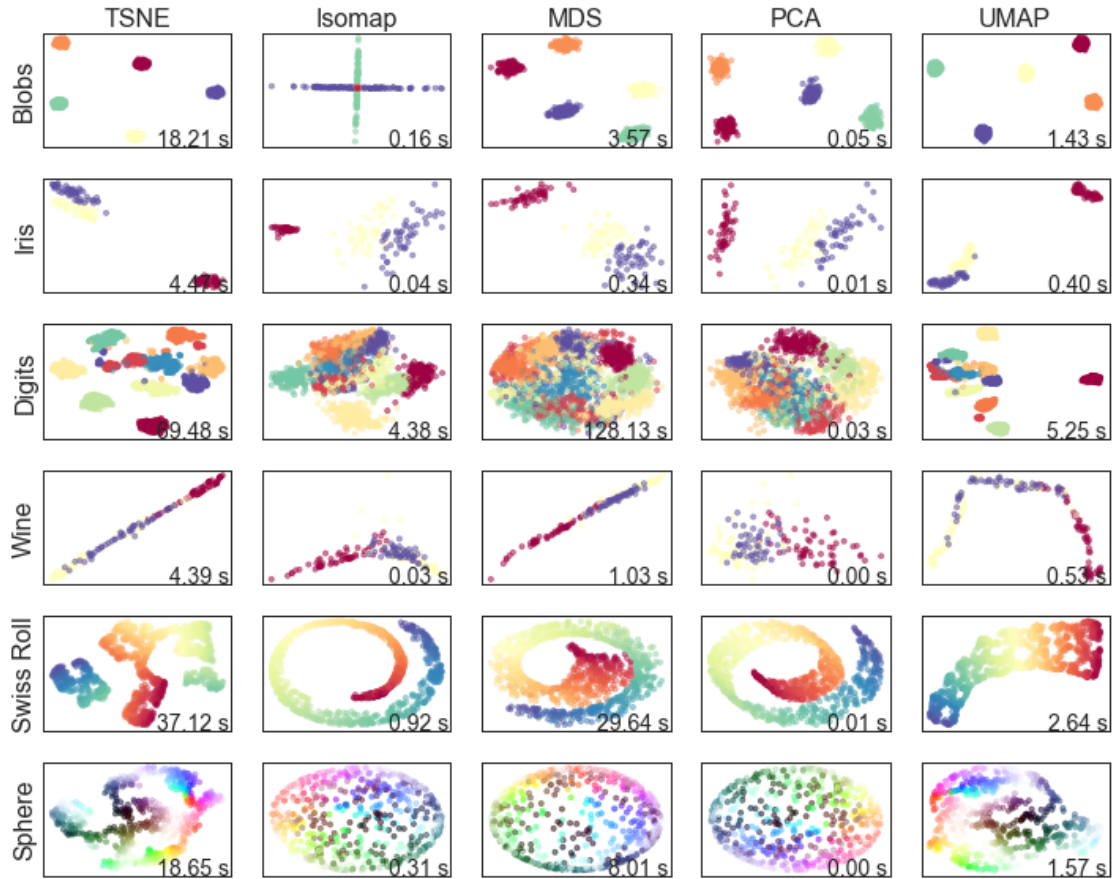
    n_rows = len(test_data)
    n_cols = len(reducers)
    ax_index = 1
    ax_list = []

    # plt.figure(figsize=(9 * 2 + 3, 12.5))
    plt.figure(figsize=(10, 8))
    plt.subplots_adjust(
        left=.02, right=.98, bottom=.001, top=.96, wspace=.05, hspace=.01
    )
    for data, labels in test_data:
        for reducer, args in reducers:
            start_time = time.time()
            embedding = reducer(n_components=2, **args).fit_transform(data)
            elapsed_time = time.time() - start_time
            ax = plt.subplot(n_rows, n_cols, ax_index)
            if isinstance(labels[0], tuple):
                ax.scatter(*embedding.T, s=10, c=labels, alpha=0.5)
            else:
                ax.scatter(
                    *embedding.T, s=10, c=labels, cmap="Spectral", alpha=0.5
                )
            ax.text(
                0.99,
                0.01,
                "{:.2f} s".format(elapsed_time),
                transform=ax.transAxes,
                size=14,
                horizontalalignment="right",
            )
            ax_list.append(ax)
            ax_index += 1
    plt.setp(ax_list, xticks=[], yticks=[])

    for i in np.arange(n_rows) * n_cols:
        ax_list[i].set_ylabel(dataset_names[i // n_cols], size=16)
    for i in range(n_cols):
        ax_list[i].set_xlabel(repr(reducers[i][0]()).split("(")[0], size=16)
        ax_list[i].xaxis.set_label_position("top")

    plt.tight_layout()
    plt.show()

```



1.0.5 UMAP on word embeddings of NIPS abstracts

```
In [32]: df_papers = pd.read_csv("../data/nips-papers/papers.csv") # papers
df_papers_clean=pd.ExcelFile("../data/abstract2.xlsx").parse('Sheet1') #cleaned version

#change Nulls to empty strings
df_papers_clean.abstract2 = df_papers_clean.abstract2.fillna('')

#convert to lower case
df_papers_clean.abstract2 = df_papers_clean.abstract2.str.lower()

#join year
df_papers_clean=df_papers_clean.join(df_papers.iloc[:,2].set_index('id'),on='id')

In [35]: #Turn text and abstracts into lists, rmv non alphabetic
#also define stop words
abstracts_list = df_papers_clean['abstract2'].tolist()
abstracts_list = [re.sub('[^a-zA-Z]', ' ',a) for a in abstracts_list]
```

```

# add to stop words
stops = set(stopwords.words("english"))
stops = stops.union(['setting', 'results', 'using', 'approach', 'problems', 'based', 'x', 'x'])

In [44]: # tokenize abstracts
tokenizer = RegexpTokenizer(r'\w+')

def tokenize(raw_text):
    raw = raw_text
    tokens = tokenizer.tokenize(raw)
    # Remove stopwords and words of length<4
    stopped_tokens = [i for i in tokens if not i in stops and len(i) > 3]
    return stopped_tokens

abstract_tokens=[tokenize(i) for i in abstracts_list]

In [39]: abstract_w2v = word2vec.Word2Vec(abstract_tokens, size=100, window=10, min_count=100,

In [42]: tokens = []
labels = []
for word in sorted(abstract_w2v.wv.vocab):
    tokens.append(abstract_w2v.wv[word])
    labels.append(word)

In [59]: umap_nips_reducer = umap.UMAP(metric='cosine')

umap_nips = umap_nips_reducer.fit_transform(tokens)


print(umap_nips.shape)
umap_nips[:5]

(1152, 2)

Out [59]: array([[ 1.4886876 ,  2.3720775 ],
 [ 0.6506124 ,  3.496596  ],
 [-0.05195517,  0.9964709 ],
 [-2.8000128 ,  0.23083588],
 [ 2.776519 ,  0.5102816 ]], dtype=float32)

In [65]: plt.figure(figsize=(15,8))
plt.scatter(umap_nips[:, 0], umap_nips[:, 1],s=5)
plt.gca().set_aspect('equal', 'datalim')
plt.title('UMAP projection of the NIPS dataset', fontsize=24);

```



In []: